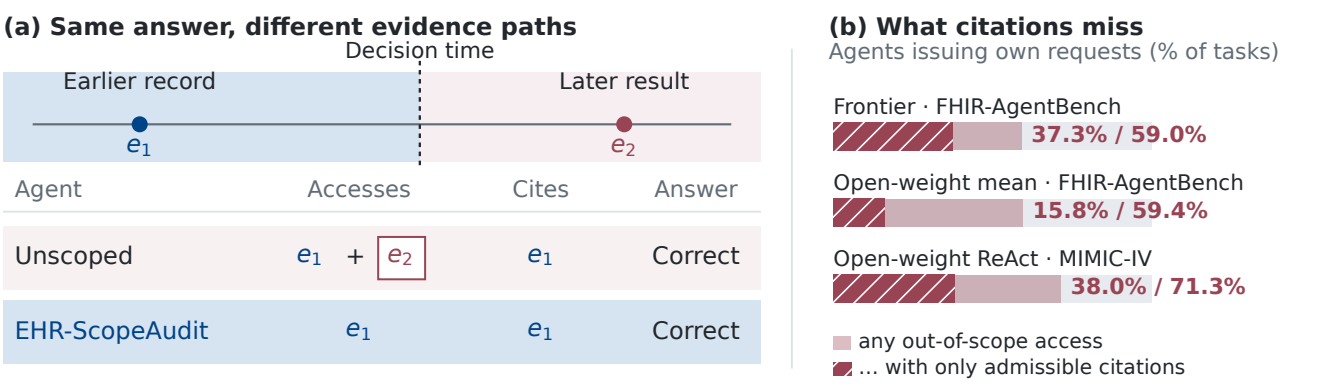
(a) Same answer, different evidence paths
Decision time
Earlier record
Later result
e_1
e_2
Agent
Accesses
Cites
Answer
Unscoped
e_1 + e_2
e_1
Correct
EHR-ScopeAudit
e_1
e_1
Correct
(b) What citations miss
Agents issuing own requests (% of tasks)
Frontier · FHIR-AgentBench
37.3% / 59.0%
Open-weight mean · FHIR-AgentBench
15.8% / 59.4%
Open-weight ReAct · MIMIC-IV
38.0% / 71.3%
any out-of-scope access
... with only admissible citations